



# MAQAO

## Performance Analysis and Optimization Framework

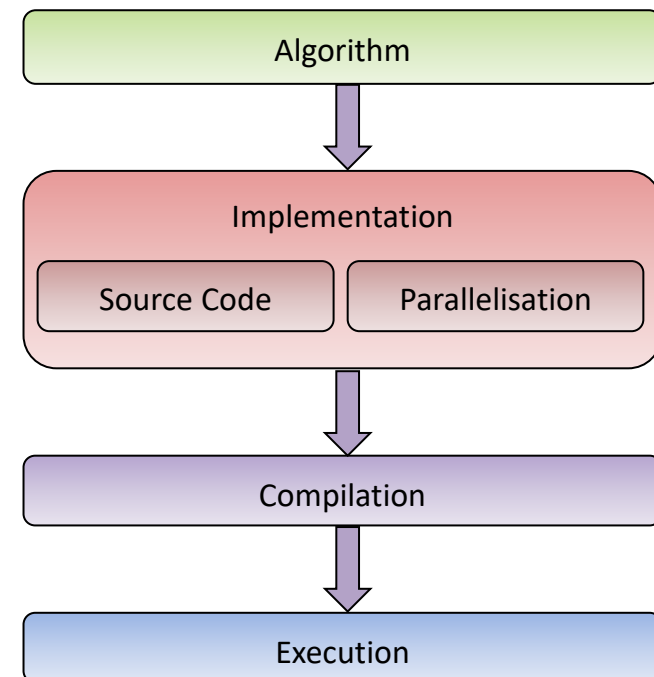
Performance Evaluation Team, University of  
Versailles

<http://maqao.exascale-computing.eu>



# Performance Analysis and Optimisation

- **Where** is the application spending most execution time and resources?
- **Why** is the application spending time there?
  - Algorithm, implementation, runtime or hardware?
  - Data access or computation?
- **How much** of an application can be optimised?
  - What would the effort/gain ratio be?
- **How** to improve the situation?
  - At which step(s) of the design process?
  - What additional information is needed?





# A Multifaceted Problem

- **Pinpointing** the performance bottlenecks
- **Identifying** the dominant issues
  - Algorithms, implementation, parallelisation, ...
- Making the **best use** of the machine features
  - Complex multicore and manycore CPUs
  - Complex memory hierarchy
- Finding the **most rewarding** issues to be fixed
  - **40%** total time, expected **10%** speedup
    - ➔ TOTAL IMPACT: **4%** speedup
  - **20%** total time, expected **50%** speedup
    - ➔ TOTAL IMPACT: **10%** speedup

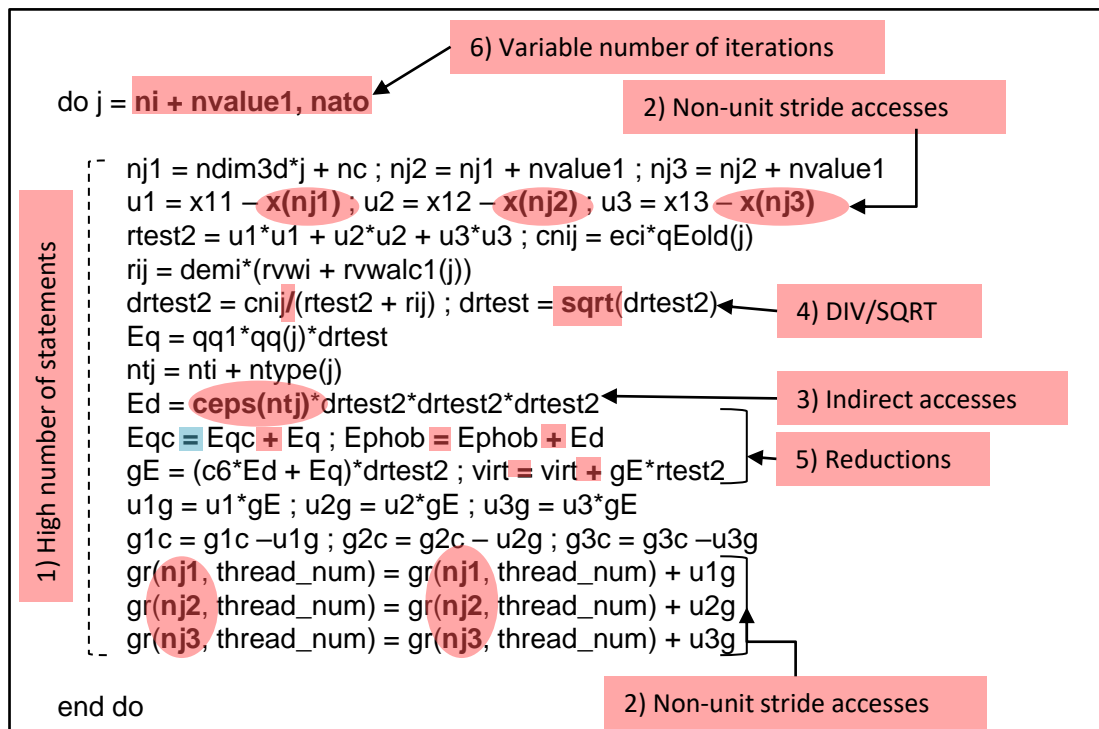


=> **Need for dedicated and complementary tools**



# Motivating Example

- Code of a loop representing ~10% walltime



Source code and associated issues:

- 1) High number of statements
- 2) Non-unit stride accesses
- 3) Indirect accesses
- 4) DIV/SQRT
- 5) Reductions
- 6) Variable number of iterations



# MAQAO: Modular Assembly Quality Analyzer and Optimizer

- Objectives:
  - Characterizing performance of HPC applications
  - **Guiding users** through optimization process
  - Estimating return of investment (**R.O.I.**)
- Characteristics:
  - **Modular tool** offering complementary views
  - Support for **Intel x86-64**, **Xeon Phi** and **AArch64** (beta version)
  - LGPL3 Open Source software
  - Developed at UVSQ since 2004
  - Binary release available as **static executable**








# Website & resources

- MAQAO website: [maqao.exascale-computing.eu](http://maqao.exascale-computing.eu)
  - Mirror: [www.maqao.org](http://www.maqao.org)
- Documentation: [maqao.exascale-computing.eu/documentation.html](http://maqao.exascale-computing.eu/documentation.html)
  - Tutorials for ONE View, LProf and CQA
  - Lua API documentation
- Latest release: [maqao.exascale-computing.eu/downloads.html](http://maqao.exascale-computing.eu/downloads.html)
  - Binary releases (2-3 per year)
  - Core sources
- Publications: [maqao.exascale-computing.eu/publications.html](http://maqao.exascale-computing.eu/publications.html)



# Success stories: Optimisation of Industrial and Academic HPC Applications

- QMC=CHEM (IRSAMC)
  - Quantum chemistry
  - Speedup: > 3x
    - Moved invocation of function with identical parameters out of loop body
- Yales2 (CORIA)
  - Computational fluid dynamics
  - Speedup: up to 2,8x
    - Removed double structure indirections
- Polaris (CEA)
  - Molecular dynamics
  - Speedup: 1,5x – 1,7x
    - Enforced loop vectorisation through compiler directives
- AVBP (CERFACS)
  - Computational fluid dynamics
  - Speedup: 1,08x – 1,17x
    - Replaced division with multiplication by reciprocal
    - Complete unrolling of loops with small number of iterations
- Ongoing effort
  - TRES CoE project codes
  - CEA DAM codes





# Partnerships

- MAQAO was funded by UVSQ, Intel (2005-2020) and CEA (French department of energy) through Exascale Computing Research (ECR) and the French Ministry of Industry through various FUI/ITEA projects (H4H, COLOC, PerfCloud, ELCI, MB3, etc...)



- Provides core technology to be integrated with other tools:
  - TAU performance tools with MADRAS patcher through MIL (MAQAO Instrumentation Language)
  - ATOS bullxprof with MADRAS through MIL
  - Intel Advisor
  - INRIA Bordeaux HWLOC
- PeXL ISV also contributes to MAQAO:
  - Commercial performance optimization expertise
  - Training and software development
  - [www.pexl.eu](http://www.pexl.eu)







# MAQAO Team and Collaborators

- **MAQAO Team**

- William Jalby, Prof.
- Cédric Valensi, Ph.D.
- Emmanuel Oseret, Ph.D.
- Mathieu Tribalat, M.Sc.Eng
- Salah Ibn Amar, M.Sc.Eng
- Hugo Bolloré , M.Sc.Eng
- Kévin Camus, Eng.
- Aurélien Delval, Eng.
- Max Hoffer, Eng.

- **Collaborators**

- David J. Kuck, Prof. (Intel US)
- Andrés S. Charif-Rubial, Ph.D. (start-up)
- Eric Petit, Ph.D. (Intel US)
- Pablo de Oliveira, Ph.D. (UVSQ)
- David Wong, Ph.D. (Intel US)
- Othman Bouizi, Ph.D. (Intel US)

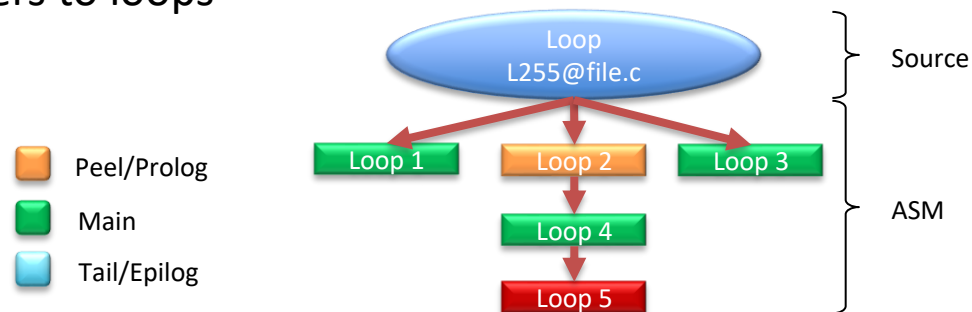
- **Past Collaborators or Team members**

- Denis Barthou, Prof. (Univ. Bordeaux)
- Jean-Thomas Acquaviva, Ph.D. (DDN)
- Stéphane Zuckerman, Ph.D. (M. Conf ENSEA)
- Julien Jaeger, Ph.D. (CEA DAM)
- Souad Koliaï, Ph.D. (CELOXICA)
- Zakaria Bendifallah, Ph.D. (ATOS)
- Tipp Moseley, Ph.D. (Google)
- Jean-Christophe Beyler, Ph.D. (Google)
- Jean-Baptiste Le Reste , M.Sc.Eng (start-up)
- Sylvain Henry, Ph.D. (start-up)
- José Noudohouennou, Ph.D. (AMD)
- Aleksandre Vardoshvili , M.Sc.Eng
- Romain Pillot, Eng
- Youenn Lebras, Ph.D. (start-up)



# Analysis at Binary Level

- Advantages of binary analysis: **What You Analyse Is What You Run**
- Issues binary analysis addresses:
  - Compiler optimizations increase the distance between the executed code and the source code
  - Source code instrumentation may prevent the compiler from applying certain transformations
- Main steps:
  - Construct high level structures (CFG, DDG, SSA, ...)
  - Relate the analyses to source code using debug information
    - A single source loop can be compiled as multiple assembly loops
    - Affecting unique identifiers to loops



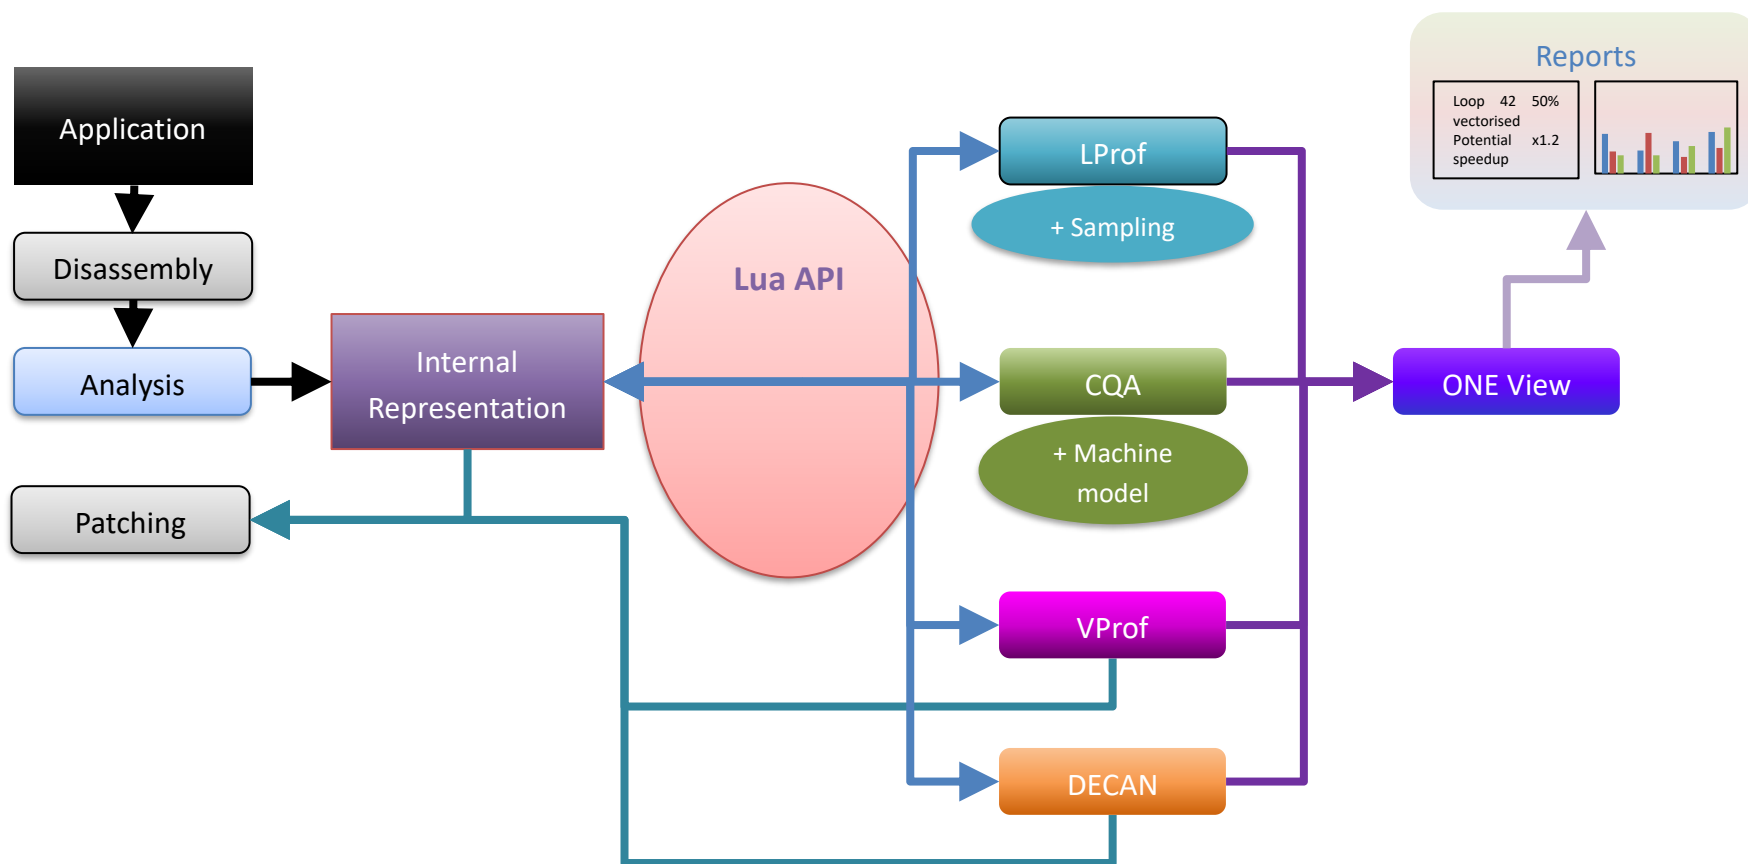


# MAQAO Main Features

- Binary layer
  - Builds internal representation from binary
  - Allows patching through binary rewriting
- Profiling
  - LProf: Lightweight sampling-based Profiler
  - VProf: Instrumentation-based Value Profiler
- Static analysis
  - CQA (Code Quality Analyzer): Evaluates the quality of the binary code and offers hints for improving it
  - UFS (Uops Flow Simulator): Cycle-accurate CPU engine simulator
- Dynamic analysis
  - DECAN (DECremental Analyzer): Modifies the application to evaluate the impact of groups of instructions on performance
- Performance view aggregation module
  - ONE View: Invokes the modules and produces reports aggregating their results




# MAQAO Main Structure





# MAQAO Methodology


- Decision tree



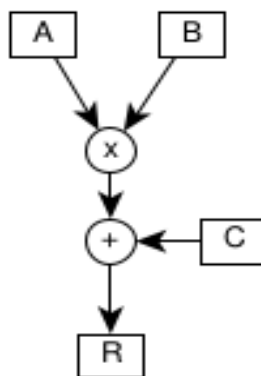


# SIMD/Vectorization/Data Parallelism

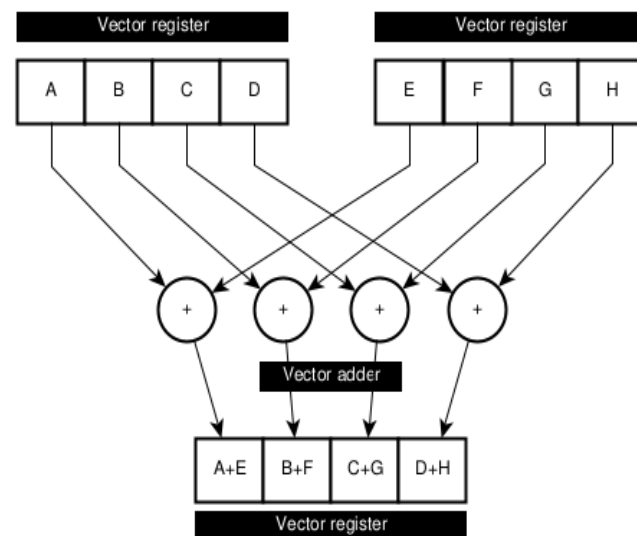
- Scalar pattern (C):  $a[i] = b[i] + c[i]$
- Vector pattern (FORTRAN):  $a(i, i + 8) = b(i, i + 8) + c(i, i + 8)$
- Benefits : increases memory bandwidth and **IPC**
- Implementations:
  - x86 : SSE, AVX, AVX512
  - ARM : Neon, SVE
- FMA/MAC: (the core operation of LinAlg/DSP algorithms)
  - Fused-Multiply-Add
  - Multiply-Accumulate



Scalar addition



FMA / MAC



Vector addition




# Compiler optimisations

- Compiler flags:
  - Loop unrolling: `-funroll-loops`
    - Reduce branches
    - Fill the pipeline (more instructions per iteration)
    - Increases memory bandwidth and IPC
  - Function inlining: `-finline-functions`
  - Vectorization: `-ftree-vectorize`, `-ftree-slp-vectorize`, ...
  - Target micro-architectures: `-march` or `-mtune` or `-xHOST`
- Compiler directives:
  - OpenMP directives: `#pragma omp simd`, `#pragma omp parallel for`, ...
  - Intel compiler specific: `#pragma simd`, `#pragma unroll`, `#pragma inline`, ...
- Compiler/language keywords/features:
  - Using `restrict` for pointers aliasing in C/C++
  - Using `inline` for function inlining in C
  - Using array sections in FORTRAN



# Memory and caches

- Computations are, in general, faster than memory accesses
- Alignment/Contiguity of memory (x86) : `posix_memalign`, `aligned_alloc`, ...
- Are caches (L1, L2, L3) used properly?
- Memory performance → Maximum bandwidth



Aligned memory access



Crossing cacheline boundary

Unaligned memory access





# MAQAO LProf: Lightweight Profiler

- **Goal:** Localization of application hotspots
- **Features:**
  - Lightweight
  - **Sampling** based
  - Access to hardware counters
  - Analysis at function and loop granularity
- **Strengths:**
  - **Non intrusive:** No recompilation necessary
  - **Low overhead**
  - Agnostic with regard to parallel runtime

MAQAO Global Application Functions Loops Help						
Functions and Loops						
Right-click on a line to display the associated load balancing. Double click on a loop to display its analysis details.						
Name	Module	Coverage (%)	Time (s)	Nb Threads	Deviation	
o binvrhs	bt-mz.C.16	23.19	13.66	64	1.73	
▼ y_solve	bt-mz.C.16	13.09	7.71	64	1.08	
▼ Loop 204 - y_solve.f:53-407 - bt-mz.C.16		12.84	7.56			
▼ Loop 205 - y_solve.f:54-407 - bt-mz.C.16		12.84	7.56			
▼ Loop 207 - y_solve.f:54-398 - bt-mz.C.16		12.84	7.56			
o Loop 211 - y_solve.f:145-307 - bt-mz.C.16		7.06	4.16			
o Loop 213 - y_solve.f:55-137 - bt-mz.C.16		4.43	2.61			
o Loop 206 - y_solve.f:394-398 - bt-mz.C.16		0.88	0.52			
o Loop 209 - y_solve.f:337-360 - bt-mz.C.16		0.33	0.19			
o Loop 210 - y_solve.f:145-307 - bt-mz.C.16		0.09	0.05			
o Loop 212 - y_solve.f:55-137 - bt-mz.C.16		0.05	0.03			
► x_solve	bt-mz.C.16	12.49	7.35	64	1.02	
o _INTERNAL_25_...src_kmp_barrier.cpp:ce635104:___kmp_hyper_barrier_release(barrier_type, kmp_info*, int, int, void*)	libiomp5.so	12.36	7.28	64	8.22	
► matmul_sub	bt-mz.C.16	11.95	7.04	64	0.92	
► z_solve	bt-mz.C.16	8.03	4.73	64	0.57	
► compute_rhs	bt-mz.C.16	7.69	4.53	64	0.59	
► matvec_sub	bt-mz.C.16	3.33	1.96	64	0.34	
o MPI_DL_CH3L_Progress	libmpi.so.12.0	1.85	1.09	16	1.91	
o binvrhs	bt-mz.C.16	0.49	0.29	64	0.05	
► lhsinit	bt-mz.C.16	0.45	0.26	64	0.04	
► add#omp_loop_0	bt-mz.C.16	0.32	0.19	64	0.03	
o system_call_after_swaps	SYSTEM CALL	0.22	0.13	63	0.12	
o _INTERNAL_25_...src_kmp_barrier.cpp:ce635104:___kmp_hyper_barrier_gather(barrier_type, kmp_info*, int, int, void (*) (void*, void*), void*)	libiomp5.so	0.16	0.1	64	0.24	
o sysret_check	SYSTEM CALL	0.14	0.08	64	0.08	
o __kmp_yield	libiomp5.so	0.13	0.07	57	0.08	
o apic_timer_interrupt	SYSTEM CALL	0.13	0.08	64	0.02	
► copy_x_face#omp_loop_0	bt-mz.C.16	0.12	0.07	64	0.03	
► exact_solution	bt-mz.C.16	0.12	0.07	64	0.01	
o update_curr	SYSTEM CALL	0.12	0.07	64	0.05	
o __audit_syscall_entry	SYSTEM CALL	0.12	0.07	64	0.08	
o __schedule	SYSTEM CALL	0.12	0.07	63	0.07	
o task_tick_fair	SYSTEM CALL	0.1	0.06	64	0.02	
► copy_y_face#omp_loop_0	bt-mz.C.16	0.1	0.06	64	0.02	
o cpuacct_charge	SYSTEM CALL	0.1	0.06	64	0.05	
o intel_pstate_update_util	SYSTEM CALL	0.1	0.06	64	0.04	
o ktime_get	SYSTEM CALL	0.09	0.05	64	0.02	



# MAQAO CQA: Code Quality Analyzer

- Goal: **Assist developers** in improving code performance
- Features:
  - Static analysis: **no execution** of the application
  - Allows **cross-analysis** of/on multiple architectures
  - Evaluate the **quality** of compiler generated code
  - Proposes **hints and workarounds** to improve quality / performance
  - **Loop centric**
    - In HPC loops cover most of the processing time
  - Targets **compute-bound** codes

**Static Reports**

**▼ CQA Report**

The loop is defined in /tmp/NPB3.3.1-MZ/NPB3.3-MZ-MPI/BT-MZ/z\_solve.f:415-423

**▼ Path 1**

2% of peak computational performance is used (0.77 out of 32.00 FLOP per cycle (GFLOPS @ 1GHz))

**gain** **potential** **hint** **expert**

**Code clean check**

Detected a slowdown caused by scalar integer instructions (typically used for address computation). By removing them, you can lower the cost of an iteration from 65.00 to 57.00 cycles (1.14x speedup).

**Workaround**

- Try to reorganize arrays of structures to structures of arrays
- Consider to permute loops (see vectorization gain report)
- To reference allocatable arrays, use "allocatable" instead of "pointer" pointers or qualify them with the "contiguous" attribute (Fortran 2008)
- For structures, limit to one indirection. For example, use a\_b%c instead of a%b%c with a\_b set to a%b before this loop

**Vectorization**

Your loop is not vectorized. 8 data elements could be processed at once in vector registers. By vectorizing your loop, you can lower the cost of an iteration from 65.00 to 8.12 cycles (8.00x speedup).

**Workaround**

- Try another compiler or update/tune your current one:
  - use the vec-report option to understand why your loop was not vectorized. If "existence of vector dependences", try the IVDEP directive. If, using IVDEP, "vectorization possible but seems inefficient", try the VECTOR ALWAYS directive.
- Remove inter-iterations dependences from your loop and make it unit-stride:
  - If your arrays have 2 or more dimensions, check whether elements are accessed contiguously and, otherwise, try to permute loops accordingly: Fortran storage order is column-major: do i do j a(i,j) = b(i,j) (slow, non stride 1) => do i do j a(j,i) = b(i,j) (fast, stride 1)
  - If your loop streams arrays of structures (AoS), try to use structures of arrays instead (SoA): do i a(i)%x = b(i)%x (slow, non stride 1) => do i a%x(i) = b%x(i) (fast, stride 1)

**Execution units bottlenecks**

Found no such bottlenecks but see expert reports for more complex bottlenecks.




# MAQAO CQA Main Concepts

- Applications only exploit at best 5% to 10% of the peak performance

- Main elements of analysis:

- Peak performance
- Execution pipeline
- Resources/Functional units

Same instruction – Same cost



Process up to  
8X (SP) data

- Key performance levers for core level efficiency:
  - Vectorisation
  - Avoiding high latency instructions if possible (e.g. DIV/SQRT)
  - Guiding the compiler code optimisation
  - Reorganizing memory and data structures layout



# MAQAO CQA Guiding the compiler and hints

- Compiler can be driven using flags, pragmas and keywords:
  - Ensuring full use of architecture capabilities (e.g. using flag -xHost on AVX capable machines)
  - Forcing optimization (unrolling, vectorization, alignment...)
  - Bypassing conservative behaviour when possible (e.g., 1/X precision)
- Hints for implementation changes
  - Improve data access patterns
    - Memory alignment
    - Loop interchange
    - Change loop stride
    - Reshaping arrays of structures
  - Avoid instructions with high latency (SQRT, DIV, GATHER, SCATTER, ...)



# MAQAO CQA Advanced Features

## Vector Efficiency

- Ex: vectorized SSE code on AVX machine
- Compiler: “LOOP WAS VECTORIZED”
- In reality 50% vectorization speedup loss
- CQA:
  - vectorization ratio: 100% (“all instructions vectorized”)
  - vec. efficiency ratio: 50% (“but using only half vector width”)
  - hint: “recompile with `-xHost`” (on Intel compilers)

128 bits


vectorized:

*Vec. ratio = 100%*

ADDPD (xmm)

MULPD (xmm)

etc...



256 bits

vectorized:

*Vec. ratio = 100%*

VADDPD (ymm)

VMULPD (ymm)

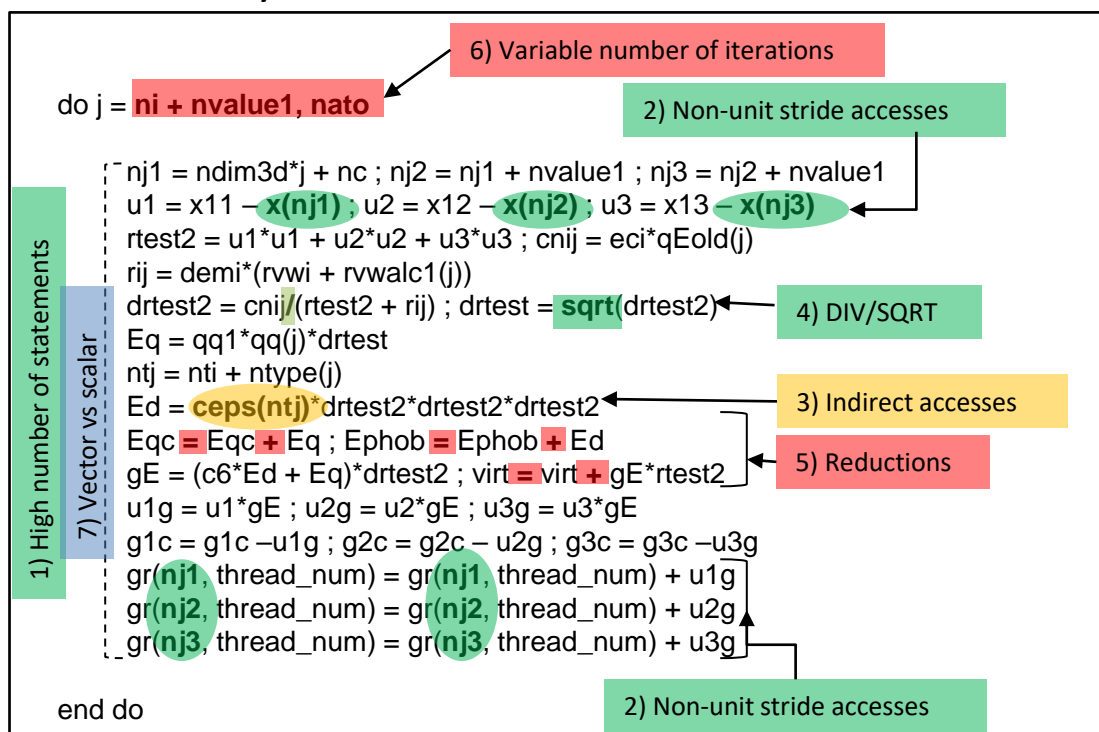
etc...





# MAQAO CQA Application to Motivating Example

## Issues identified by CQA



CQA can detect and provide hints to resolve most of the identified issues:

- 1) High number of statements
- 2) Non-unit stride accesses
- 3) Indirect accesses
- 4) DIV/SQRT
- 5) Reductions
- 6) Variable number of iterations
- 7) Vector vs scalar



# MAQAO CQA: Code Quality Analyzer

## Application to motivating example

**Gain** **Potential gain** **Hints** **Experts only**

### Vectorization

Your loop is partially vectorized.  
Only 28% of vector register length is used (average across all SSE/AVX instructions).  
By fully vectorizing your loop, you can lower the cost of an iteration from 57.00 to 21.50 cycles (2.65x speedup).  
51% of SSE/AVX instructions are used in vector version (process two or more data elements in vector registers):

- 24% of SSE/AVX loads are used in vector version.
- 0% of SSE/AVX stores are used in vector version.

Since your execution units are vector units, only a fully vectorized loop can use their full power.

**Proposed solution(s):**

- Try another compiler or update/tune your current one:
  - use the vec-report option to understand why your loop was not vectorized. If "existence of vector dependences", try the IVDEP directive. If, using IVDEP, "vectorization possible but seems inefficient", try the VECTOR ALWAYS directive.
- Remove inter-iterations dependences from your loop and make it unit-stride:
  - If your arrays have 2 or more dimensions, check whether elements are accessed contiguously and, otherwise, try to permute loops accordingly:  
Fortran storage order is column-major: do i do j a(i,j) = b(i,j) (slow, non stride 1) => do i do j a(j,i) = b(i,j) (fast, stride 1)
  - If your loop streams arrays of structures (AoS), try to use structures of arrays instead (SoA):  
do i a(i)%x = b(i)%x (slow, non stride 1) => do i a%x(i) = b%x(i) (fast, stride 1)

### Execution units bottlenecks

Performance is limited by:

- execution of divide and square root operations (the divide/square root unit is a bottleneck)
- execution of INT/FP operations in vector registers (the VPU is a bottleneck)

By removing all these bottlenecks, you can lower the cost of an iteration from 57.00 to 48.00 cycles (1.19x speedup).

**Proposed solution(s):**

- Reduce the number of division or square root instructions.  
If denominator is constant over iterations, use reciprocal (replace x/y with x\*(1/y)). Check precision impact. This will be done by your compiler with no-prec-div or Ofast.  
Check whether you really need double precision. If not, switch to single precision to speedup execution.
- Reduce arithmetical operations on array elements

**Gain** **Potential gain** **Hints** **Experts only**

### FMA

Detected 48 FMA (fused multiply-add) operations.  
Presence of both ADD/SUB and MUL operations.

**Proposed solution(s):**

Try to change order in which elements are evaluated (using parentheses)  
MUL operations to enable your compiler to generate FMA instructions where possible.  
For instance a + b\*c is a valid FMA (MUL then ADD).  
However (a+b)\*c cannot be translated into an FMA (ADD then MUL).

**Gain** **Potential gain** **Hints** **Experts only**

### Slow data structures

Detected data structures (typically arrays) that cannot be efficiently read/written

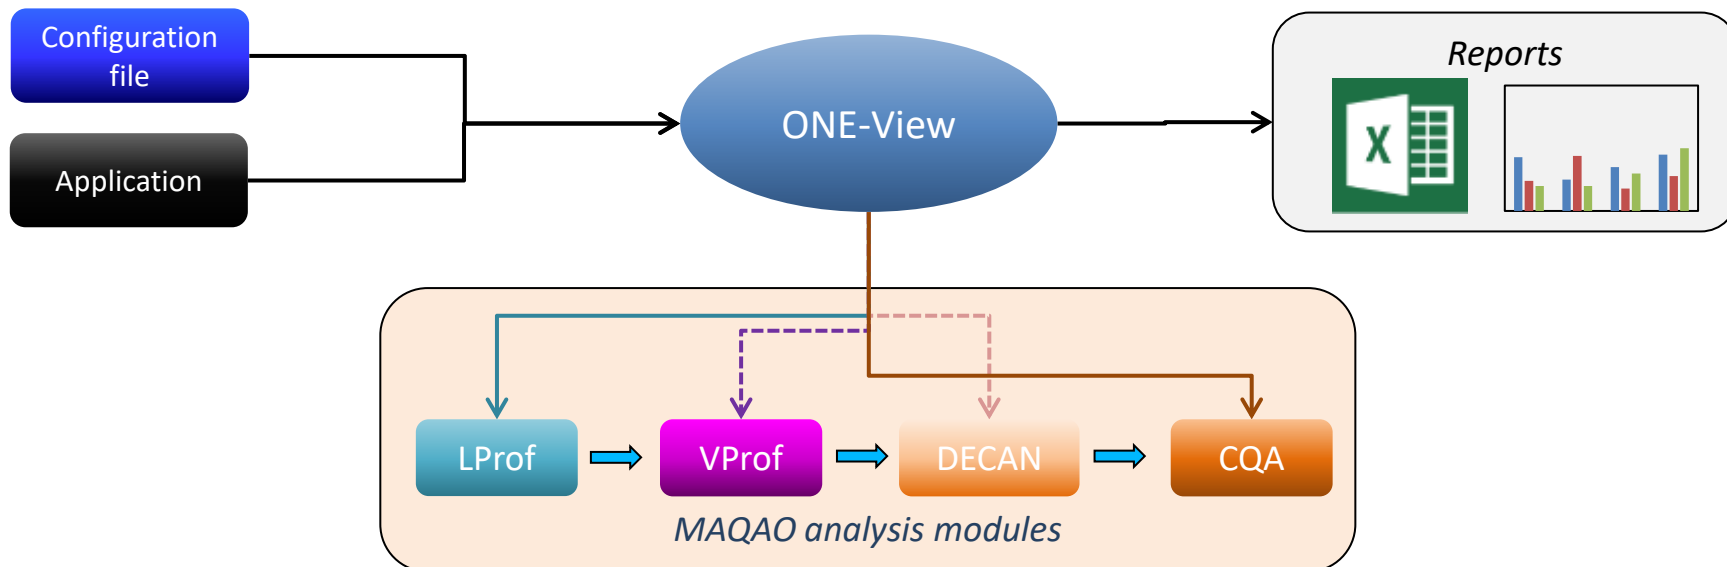
- Constant non-unit stride: 1 occurrence(s)
- Irregular (variable stride) or indirect: 1 occurrence(s)

- 1) High number of statements
- 2) Non-unit stride accesses
- 3) Indirect accesses
- 4) DIV/SQRT
- 5) Reductions
- 6) Variable number of iterations
- 7) Vector vs scalar



# MAQAO ONE View: Performance View Aggregator

- Goal: **Automating** the whole analysis process
  - Invoke multiple MAQAO modules
  - Generate **aggregated performance views**
    - Reports in HTML or XLS format

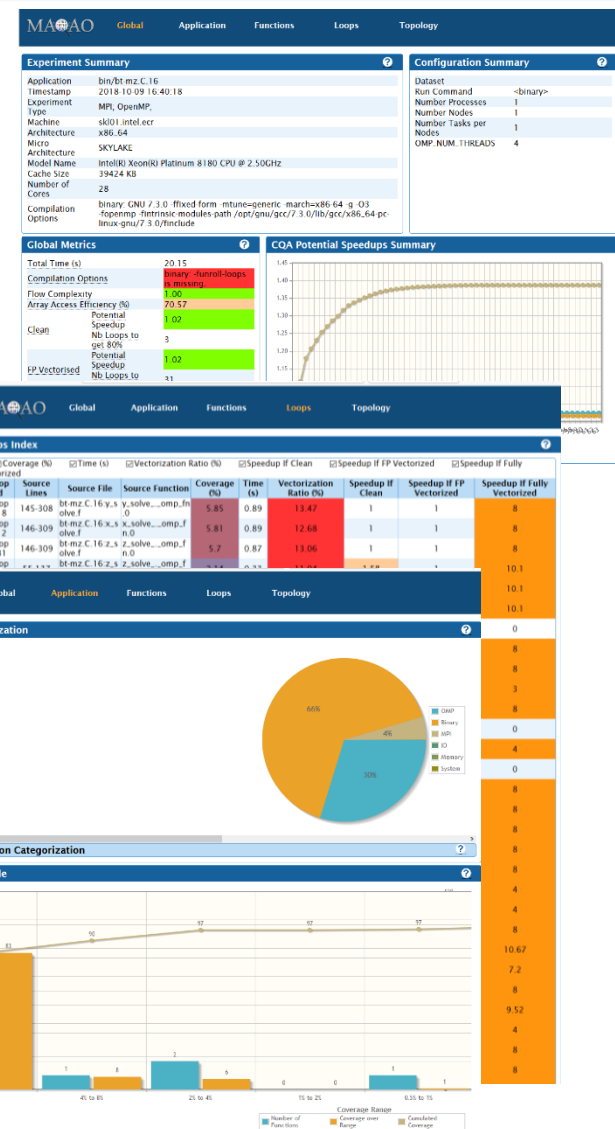






# MAQAO ONE View: Performance View Aggregator

- Main steps:
  - Invokes LProf to **identify hotspots**
  - Invokes CQA and other modules on **loop hotspots**
- Available results:
  - **Speedup** predictions
  - Global **code quality** metrics
  - **Hints** for improving performance
  - Detailed analyses results
  - Parallel **efficiency** analysis





# ONE View Reports Levels

- ONE VIEW ONE
  - Requires a single run of the application
  - Profiling of the application using LProf
  - Static analysis using CQA
- ONE VIEW TWO (includes analyses from report ONE)
  - Requires 3 or 4 runs on average
  - Value profiling using VProf to identify loop iteration count
  - Decremental analysis for L1 projection using DECAN
- ONE VIEW THREE (includes analyses from report TWO)
  - Requires 20 to 30 runs
  - Decremental analyses using all DECAN variants
  - Collects hardware performance events
- Comparison mode
  - Comparison of multiple runs (iso-binary or iso-source)
  - Allows to evaluate scalability or compare performance across different datasets, compilers, or hardware platforms



# Analysing an application with MAQAO

- ONE View execution
- Provide all parameters necessary for executing the application
  - Parameters can be passed on the command line or as a configuration file

```
$ maqao oneview -R1 ./myexe
```

```
$ maqao oneview --create-report=one --executable=./myexe --mpi_command="mpirun -n 16"
```

```
$ maqao oneview --create-report=one --config=my_config.lua
```

- Analyses can be tweaked if necessary
- ONE View can reuse an existing experiment directory to perform further analyses
- Results available in HTML format by default
  - XLS spreadsheets and textual output generation are also available
- Online help is available:

```
$ maqao oneview --help
```



# Analysing an application with MAQAO

MAQAO modules can be invoked separately for advanced analyses

- LProf
  - Profiling

```
$ maqao lprof xp=exp_dir --mpi-command="mpirun -n 16" -- ./myexe
```

- Display functions profile

```
$ maqao lprof xp=exp_dir -df
```

- Displaying the results from a ONE View run

```
$ maqao lprof xp=oneview_xp_dir/lprof_npsu -df
```

- CQA

```
$ maqao cqa loop=42 myexe
```

Online help is available:

```
$ maqao lprof --help
```

```
$ maqao cqa --help
```



Thanks for your attention

**QUESTIONS ?**



# NAVIGATING ONE VIEW REPORTS



# MAQAO ONE View Global Summary

- Experiment summary
  - Characteristics of the machine where the experiment took place
- Global metrics
  - General quality metrics derived from MAQAO analyses
  - Global speedup predictions
    - Speedup prediction depending on the number of vectorised loops
    - Ordered speedups to identify the loops to optimise in priority





# ONE View Global Metrics

- Global metrics
  - General quality metrics derived from MAQAO analyses
  - Global speedup predictions
- Potential speedups
  - Speedup prediction depending on the number of optimised loops
  - Ordered speedups to identify the loops to optimise in priority
- $Global\ Speedup = \sum_{loops} coverage * potential\ speedup$
- LProf provides coverage of the loops
- CQA and DECAN provide speedup estimation for loops
  - Speedup if loop vectorised or without address computation
  - All data in L1 cache





# MAQAO ONE View: Functions Profiling

## Identifying hotspots

- Exclusive coverage
- Load balancing across threads
- Loops nests by functions

### ▼ matmul\_sub

○ Loop 230 - solve\_subs.f:71-175 - bt-mz.C.16

○ Loop 231 - solve\_subs.f:71-175 - bt-mz.C.16

### ▼ z\_solve

▼ Loop 232 - z\_solve.f:53-423 - bt-mz.C.16

▼ Loop 233 - z\_solve.f:54-423 - bt-mz.C.16

▼ Loop 236 - z\_solve.f:54-423 - bt-mz.C.16

○ Loop 239 - z\_solve.f:146-308 - bt-mz.C.16

○ Loop 235 - z\_solve.f:55-137 - bt-mz.C.16

○ Loop 234 - z\_solve.f:415-423 - bt-mz.C.16

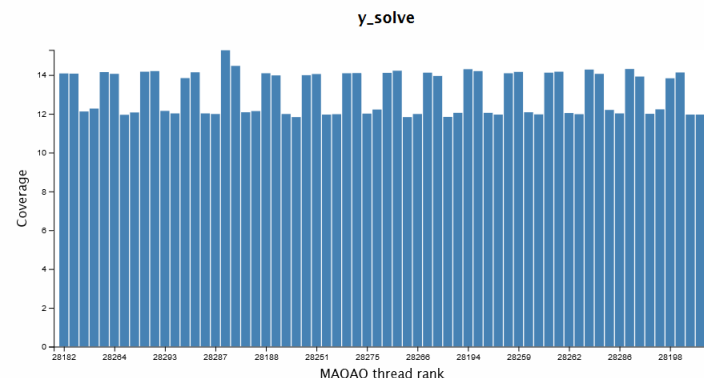

Single

Outermost

Inbetween

Inbetween

Innermost





# MAQAO ONE View Loop Profiling Summary

- Identifying loop hotspots
- Vectorisation information
- Potential speedups by optimisation
  - Clean: Removing address computations
  - FP Vectorised: Vectorising floating-point computations
  - Fully Vectorised: Vectorising floating-point computations and memory accesses

MAQAO Global Application Functions <b>Loops</b> Topology											
Show Innermost Profile Open Expert Summary											
Loops Index ?											
► Filters ?											
<input checked="" type="checkbox"/> Coverage (%) <input checked="" type="checkbox"/> Level <input checked="" type="checkbox"/> Time (s) <input checked="" type="checkbox"/> Vectorization Ratio (%) <input checked="" type="checkbox"/> Speedup If Clean <input checked="" type="checkbox"/> Speedup If FP Vectorized <input checked="" type="checkbox"/> Speedup If Fully Vectorized <input checked="" type="checkbox"/> Speedup If Data in L1 <input type="checkbox"/> Select none <input type="checkbox"/> Select All Speed-Ups <input type="checkbox"/> Select All Efficiencies											
Loop id	Source Location	Source Function	Coverage (%)	Level	Time (s)	Vectorization Ratio (%)	Speedup If Clean	Speedup If FP Vectorized	Speedup If Fully Vectorized	Speedup If Data in L1	
18403	qmcpack::MultiBsplineValue.h pp:56-57	qmcplusplus::BsplineSet >::evaluate	26.71	Innermost	3.61	100	1	1	1	8.25	
26027	qmcpack::cmath:261-464	qmcplusplus::SoaDistanceTableAA::moveOnSphere	12.01	Single	1.62	100	1	1	1	1.03	
18424	qmcpack::MultiBsplineVGLH.h pp:187-207	qmcplusplus::BsplineSet >::evaluate	10.81	Innermost	1.46	100	1.06	1	1	4.15	
18474	qmcpack::MultiBsplineVGLH.h pp:187-207	qmcplusplus::BsplineSet >::evaluate_notranspose	4.84	Innermost	0.65	100	1.06	1	1	4.52	
26026	qmcpack::cmath:261-464	qmcplusplus::SoaDistanceTableAA::evaluate	2.78	Single	0.38	100	1	1	1	1.05	
26028	qmcpack::cmath:261-464	qmcplusplus::SoaDistanceTableAA::move	2.64	Single	0.36	100	1	1	1	1.03	
8754	qmcpack::CoulombPBCAA.cpp:425-427	qmcplusplus::CoulombPBCAA::evalSR	1.57	Innermost	0.21	0	1.64	2.59	7.67	1.01	
12711	qmcpack::BsplineFunction.h:690-695	qmcplusplus::J2OrbitalSoA >::ratioGrad	1.41	Innermost	0.19	0	1.3	1	16	1.08	
18501	qmcpack::SplineC2RAdaptor.h:325-373	void qmcplusplus::SplineC2RSoA::assign_vgl >, qmcplusplus::Vector, std::allocator > > >	1.22	Single	0.16	100	1.01	1	1	3.51	



# MAQAO ONE View Scalability Reports

## Application View



- Coverage per category
  - Comparison of categories for each run

- Coverage per parallel efficiency

- $$Efficiency = \frac{T_{sequential}}{T_{parallel} * N_{threads}}$$

- Distinguishing functions only represented in parallel or sequential

- Displays efficiency by coverage





# BACKUP SLIDES



# MAQAO History

- 2004: Begun development
  - Focusing on Intel Itanium architecture
  - Analysis of assembly files
- 2006: Transition to Intel x86-64
- 2009: Binary analysis support
  - First version of decremental analysis
- 2012: Support of KNC architecture
- 2014: Profiling features
- 2015: First version of ONE View
- 2017: Prototype support of ARM architecture
- 2018: Scalability mode
- 2020: Comparison mode
- 2022: Support of ARM (beta)

The screenshot displays the MAQAO Web Interface. On the left, the 'Functions' panel shows the project name 'L\_rbgauiss\_50\_tree\_reduce2\_2.0' and the selected loop 'Loop SRC L30'. The 'Loop's DDG' panel shows a graph with nodes representing instructions and edges representing data dependencies. On the right, the 'Run' panel shows the assembly code for the selected loop, including instructions like 'float rsum;', 'float (\*array)[800000];', and 'float urel;'. The code is annotated with comments and macros.

